

NGBoost: Natural Gradient Boosting for Probabilistic Prediction

Tony Duan^{*1} Anand Avati^{*1} Daisy Yi Ding¹ Khanh K. Thai² Sanjay Basu³ Andrew Ng¹
Alejandro Schuler²

Abstract

We present Natural Gradient Boosting (NGBoost), an algorithm for generic probabilistic prediction via gradient boosting. Typical regression models return a point estimate, conditional on covariates, but probabilistic regression models output a full probability distribution over the outcome space, conditional on the covariates. This allows for predictive uncertainty estimation — crucial in applications like healthcare and weather forecasting. NGBoost generalizes gradient boosting to probabilistic regression by treating the parameters of the conditional distribution as targets for a multiparameter boosting algorithm. Furthermore, we show how the *Natural Gradient* is required to correct the training dynamics of our multiparameter boosting approach. NGBoost can be used with any base learner, any family of distributions with continuous parameters, and any scoring rule. NGBoost matches or exceeds the performance of existing methods for probabilistic prediction while offering additional benefits in flexibility, scalability, and usability.

1. Introduction

Many important supervised machine learning problems are regression problems. Weather forecasting (predicting temperature of the next day based on today’s atmospheric variables (Gneiting and Katzfuss, 2014)) and clinical prediction (predicting time to mortality with survival prediction on structured medical records of the patient (Avati et al., 2018)) are important examples.

Most machine learning methods tackle this problem with

^{*}Equal contribution ¹Stanford University, Stanford, California, United States ²Kaiser Permanente Northern California Division of Research, Oakland, California, United States ³Harvard Medical School, Cambridge, Massachusetts, United States. Correspondence to: Anand Avati <avati@cs.stanford.edu>, Alejandro Schuler <alejandro.schuler@gmail.com>.




Figure 1. Prediction intervals for a toy 1-dimensional probabilistic regression problem, fit via NGBoost. The dots represent data points. The thick black line is the predicted mean after fitting the model. The thin gray lines are the upper and lower quantiles covering 95% of the prediction distribution.

point prediction, returning a single “best guess” prediction (e.g. the temperature tomorrow will be 16C). However, in these fields it is often important to be able to quantify uncertainty in the prediction or be able to answer multiple questions on the fly (e.g. what’s the probability it will be between 18 and 20C? What about <15C?) (Kruchten, 2016).

In order to answer arbitrary questions about the probability of events conditional on covariates, we must estimate the conditional probability distribution $P(y|x)$ for each value of x instead of producing a point estimate like $\mathbb{E}[y|x]$. This is called *probabilistic regression*. Probabilistic regression is increasingly being used in fields like meteorology and healthcare (Avati et al., 2019; Gneiting and Raftery, 2005; 2007).

Probabilistic estimation is already the norm in *classification* problems. Although some classifiers (e.g. standard support vector machines) only return a predicted class label, most are capable of returning estimated probabilities for each class; effectively, a conditional mass function.

However, existing methods for probabilistic *regression* are either inflexible, slow, or inaccessible to non-experts. Any mean-estimating regression method can be made probabilistic by assuming homoscedasticity and estimating an unconditional noise model, but homoscedasticity is a




Figure 2. NGBoost is modular with respect to choice of base learner, distribution, and scoring rule.

strong assumption and the process requires some statistical know-how. Generalized Additive Models for Shape, Scale, and Location (GAMLSS) allow heteroscedasticity but are restricted to a pre-specified model form (Stasinopoulos et al., 2007). Bayesian methods naturally generate predictive uncertainty estimates by integrating predictions over the posterior, but exact solutions to Bayesian models are limited to simple models, and calculating the posterior distribution of more powerful models such as Neural Networks (NN) (Neal, 1996) and Bayesian Additive Regression Trees (BART) (Chipman et al., 2010) is difficult. Inference in these models requires computationally expensive approximation via, for example, MCMC sampling. Moreover, sampling-based inference requires some statistical expertise and thus limits the ease-of-use of Bayesian methods. Among non-parametric Bayesian approaches, scalability to very large datasets is a challenge (Rasmussen and Williams, 2005). Bayesian Deep Learning is gaining popularity (Gal and Ghahramani, 2016; Lakshminarayanan et al., 2017), but, while neural networks have empirically excelled at perception tasks (such as with visual and audio input), they are not out-of-the-box models and they often perform on par with traditional methods when data are tabular. Small training datasets and informative prior specification are also challenges for Bayesian neural nets.

Meanwhile, Gradient Boosting Machines (GBMs) (Chen and Guestrin, 2016; Friedman, 2001) are a set of highly-modular methods that work out-of-the-box and perform well on structured input data, even with relatively small datasets. This can be seen in their empirical success on Kaggle and other competitions (Chen and Guestrin, 2016). In classification tasks, their predictions are probabilistic by default (by use of the sigmoid or softmax link function). But in regression tasks, they output only a scalar value. Under a squared-error loss these scalars can be interpreted as the mean of a conditional Gaussian distribution with some (unknown) constant variance. However, such probabilistic interpretations have little use if the variance is assumed constant. The predicted distribution needs to have at least two degrees of freedom (two parameters) to effectively convey both the magnitude and the uncertainty of the prediction, as illustrated in Figure 1. It is precisely this problem of simultaneous boosting of multiple parameters from the base learners which makes probabilistic forecasting with GBMs a challenge, and NGBoost addresses this

with a multiparameter boosting approach and the use of natural gradients (Amari, 1998).

Summary of Contributions

- i. We present Natural Gradient Boosting, a modular algorithm for probabilistic regression (sec 2.4) which uses multiparameter boosting and natural gradients to integrate any choice of:
 - Base learner (e.g. Regression Tree),
 - Parametric probability distribution (Normal, Laplace, etc.), and
 - Scoring rule (MLE, CRPS, etc.).
- ii. We present a generalization of the natural gradient to other scoring rules such as CRPS (sec 2.2).
- iii. We demonstrate empirically that NGBoost performs competitively relative to other models in its predictive uncertainty estimates as well as on point-estimation metrics (sec 3).

2. Natural Gradient Boosting

In standard prediction settings, the object of interest is an estimate of a scalar function like $\mathbb{E}[y|x]$, where x is a vector of observed features and y is the prediction target. In our setting we are interested in producing a probability distribution $P_\theta(y|x)$ (with CDF F_θ). Our approach is to assume $P_\theta(y|x)$ is of a specified parametric form, then estimate the p parameters $\theta \in \mathbb{R}^p$ of the distribution as functions of x .

2.1. Proper Scoring Rules

To begin with, we need a learning objective. In point prediction, the predictions are compared to the observed data with a loss function. The analogue in probabilistic regression is a *scoring rule*, which compares the estimated probability distribution to the observed data.

A *proper* scoring rule \mathcal{S} takes as input a forecasted probability distribution P and one observation y (outcome), and assigns a score $\mathcal{S}(P, y)$ to the forecast such that the true distribution of the outcomes gets the best score in expectation (Gneiting and Raftery, 2007). In mathematical notation, a scoring rule \mathcal{S} is a proper scoring rule if and only if it sat-

isfies

$$\mathbb{E}_{y \sim Q}[\mathcal{S}(Q, y)] \leq \mathbb{E}_{y \sim Q}[\mathcal{S}(P, y)] \quad \forall P, Q, \quad (1)$$

where Q represents the true distribution of outcomes y , and P is any other distribution (such as the probabilistic forecast from a model). Since we are working with parametric distributions, we can identify each distribution with its parameters and write the score as $\mathcal{S}(\theta, y)$.

The most commonly used proper scoring rule is the log score \mathcal{L} , which, when minimized, gives the MLE:

$$\mathcal{L}(\theta, y) = -\log P_\theta(y).$$

The CRPS is another proper scoring rule, which is generally considered a robust alternative to MLE (Gebetsberger et al., 2018). The CRPS (denoted \mathcal{C}) is defined as

$$\mathcal{C}(\theta, y) = \int_{-\infty}^y F_\theta(z)^2 dz + \int_y^{\infty} (1 - F_\theta(z))^2 dz.$$

2.2. The Generalized Natural Gradient

We will take a standard gradient descent approach to find the parameters that minimize the scoring rule by descending along the negative gradient of the score relative to the parameters at each point x . The (ordinary) gradient of a scoring rule \mathcal{S} over a parameterized probability distribution P_θ with parameter θ and outcome y with respect to the parameters is denoted $\nabla \mathcal{S}(\theta, y)$. It is the direction of steepest ascent, such that moving the parameters an infinitesimally small amount in that direction of the gradient (as opposed to any other direction) will increase the scoring rule the most. That is,

$$\nabla \mathcal{S}(\theta, y) \propto \lim_{\epsilon \rightarrow 0} \arg \max_{d: \|d\|=\epsilon} \mathcal{S}(\theta + d, y).$$

This gradient is *not* invariant to reparametrization. Consider reparameterizing P_θ to $P_{z(\theta)}(y)$ so $P_\theta(y \in A) = P_\psi(y \in A)$ for all events A when $\psi = z(\theta)$. If the gradient is calculated relative to θ and an infinitesimal step is taken in that direction, say from θ to $\theta + d\theta$ the resulting distribution will be different than if the gradient had been calculated relative to ψ and a step was taken from ψ to $\psi + d\psi$. In other words, $P_{\theta+d\theta}(y \in A) \neq P_{\psi+d\psi}(y \in A)$. Thus the choice of parametrization can drastically impact the training dynamics, even though the minima are unchanged.

The problem is that “distance” between two parameter values does not correspond to an appropriate “distance” between the distributions that those parameters identify. This motivates the natural gradient (denoted $\tilde{\nabla}$), which originated in information geometry (Amari, 1998).

Divergences. Every proper scoring rule induces a *divergence* that can serve as local distance metric in the space of

distributions. A proper scoring rule by definition satisfies the inequality of Eqn 1. The excess score of the right hand side over the left is the divergence induced by that scoring rule (Dawid and Musio, 2014):

$$D_{\mathcal{S}}(Q \| P) = \mathbb{E}_{y \sim Q}[\mathcal{S}(P, y)] - \mathbb{E}_{y \sim Q}[\mathcal{S}(Q, y)],$$

which is necessarily non-negative, and can be interpreted as a measure of difference from one distribution Q to another P . The MLE scoring rule induces the Kullback-Leibler divergence (KL divergence, or D_{KL}), while CRPS induces the L^2 divergence (Dawid, 2007; Machete, 2013).

The divergences D_{KL} and D_{L^2} are invariant to how Q and P are parametrized. Though divergences in general are not symmetric, for small changes of the parameters they are almost symmetric and can serve as a local distance metric. When used as such, a divergence induces a statistical manifold where each point in the manifold corresponds to a probability distribution (Dawid and Musio, 2014).

Natural Gradient. The generalized natural gradient is the direction of steepest ascent in Riemannian space, which is invariant to parametrization, and is defined as:

$$\tilde{\nabla} \mathcal{S}(\theta, y) \propto \lim_{\epsilon \rightarrow 0} \arg \max_{d: D_{\mathcal{S}}(P_\theta \| P_{\theta+d})=\epsilon} \mathcal{S}(\theta + d, y).$$

If we solve the corresponding the optimization problem, we obtain the natural gradient of the form

$$\tilde{\nabla} \mathcal{S}(\theta, y) \propto \mathcal{I}_{\mathcal{S}}(\theta)^{-1} \nabla \mathcal{S}(\theta, y)$$

where $\mathcal{I}_{\mathcal{S}}(\theta)$ is the Riemannian metric of the statistical manifold at θ , which is induced by the scoring rule \mathcal{S} . While the natural gradient was originally defined for the statistical manifold with the distance measure induced by D_{KL} (Martens, 2014), we provide a more general treatment here that applies to any divergence that corresponds to some proper scoring rule.

By choosing $\mathcal{S} = \mathcal{L}$ (i.e. MLE) and solving the above optimization problem, we get:

$$\tilde{\nabla} \mathcal{L}(\theta, y) \propto \mathcal{I}_{\mathcal{L}}(\theta)^{-1} \nabla \mathcal{L}(\theta, y)$$

where $\mathcal{I}_{\mathcal{L}}(\theta)$ is the Fisher Information carried by an observation about P_θ , which is defined as:

$$\mathcal{I}_{\mathcal{L}}(\theta) = \mathbb{E}_{y \sim P_\theta} [\nabla_\theta \mathcal{L}(\theta, y) \nabla_\theta \mathcal{L}(\theta, y)^T]$$

Similarly, by choosing $\mathcal{S} = \mathcal{C}$ (i.e. CRPS) and solving the above optimization problem, we get:

$$\tilde{\nabla} \mathcal{C}(\theta, y) \propto \mathcal{I}_{\mathcal{C}}(\theta)^{-1} \nabla \mathcal{C}(\theta, y)$$




Figure 3. Proper scoring rules and corresponding gradients for fitting a Normal distribution on samples $\sim N(0, 1)$. For each scoring rule, the landscape of the score (colors and contours) is identical, but the gradient fields (arrows) are markedly different depending on which kind of gradient is used.

where $\mathcal{I}_{\mathcal{C}}(\theta)$ is the Riemannian metric of the statistical manifold that uses D_{L^2} as the local distance measure, given by (Dawid, 2007):

$$\mathcal{I}_{\mathcal{C}}(\theta) = 2 \int_{-\infty}^{\infty} \nabla_{\theta} F_{\theta}(z) \nabla_{\theta} F_{\theta}(z)^T dz.$$

Using the natural gradient for learning the parameters makes the optimization problem invariant to parametrization and leads to more efficient and stable learning dynamics (Amari, 1998). Figure 3 shows the vector field of gradients and natural gradients for \mathcal{L} and \mathcal{C} on the parameter space of a Normal distribution parameterized by μ (mean) and $\log \sigma$ (logarithm of the standard deviation).

2.3. Gradient Boosting

Gradient boosting (Friedman, 2001) is a supervised learning technique where several weak learners (or base learners) are combined in an additive ensemble. The model is learnt sequentially, where the next base learner is fit against the training objective residual of the current ensemble. The output of the fitted base learner is then scaled by a learning rate and added into the ensemble.

The boosting framework can be generalized to any choice of base learner but most popular implementations use shallow decision trees because they work well in practice (Chen and Guestrin, 2016; Ke et al., 2017).

When fitting a decision tree to the gradient, the algorithm partitions the data into axis-aligned slices. Each slice of the partition is associated with a leaf node of the tree, and is made as homogeneous in its response variable (the gradients at that set of data points) as possible. The criterion of homogeneity is typically the sample variance. The prediction value of the leaf node (which is common to all the examples ending up in the leaf node) is then set to be the additive component to the predictions that minimizes the loss the most. This is equivalent to doing a “line search” in the functional optimization problem for each leaf node, and, for some losses, closed form solutions are available. For example, for squared error, the result of the line search will yield the sample mean of the response variables in the leaf.

We now consider adapting gradient boosting for prediction of parameters θ in the probabilistic regression context.

2.4. NGBoost: Natural Gradient Boosting

The NGBoost algorithm is a supervised learning method for probabilistic prediction that uses boosting to estimate the parameters of a conditional probability distribution $P(y|x)$ as functions of x . Here y could be one of several types ($\{\pm 1\}$, \mathbb{R} , $\{1, \dots, K\}$, \mathbb{R}_+ , \mathbb{N} , etc.) and x is a vector in \mathbb{R}^d . In our experiments we focus on real valued outputs, though all of our methods are applicable to other modalities

such as classification and survival prediction.

The algorithm has three modular components, which are chosen upfront as a configuration:

- Base learner (f),
- Parametric probability distribution (P_θ), and
- Proper scoring rule (\mathcal{S}).

Algorithm 1 NGBoost for probabilistic prediction

Data: Dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^n$.

Input: Boosting iterations M , Learning rate η , Probability distribution with parameter θ , Proper scoring rule \mathcal{S} , Base learner f .

Output: Scalings and base learners $\{\rho^{(m)}, f^{(m)}\}_{m=1}^M$.
 $\theta^{(0)} \leftarrow \arg \min_{\theta} \sum_{i=1}^n \mathcal{S}(\theta, y_i)$ {initialize to marginal}

```

for  $m \leftarrow 1, \dots, M$  do
  for  $i \leftarrow 1, \dots, n$  do
     $g_i^{(m)} \leftarrow \mathcal{I}_{\mathcal{S}} \left( \theta_i^{(m-1)} \right)^{-1} \nabla_{\theta} \mathcal{S} \left( \theta_i^{(m-1)}, y_i \right)$ 
  end
   $f^{(m)} \leftarrow \text{fit} \left( \left\{ x_i, g_i^{(m)} \right\}_{i=1}^n \right)$ 
   $\rho^{(m)} \leftarrow \arg \min_{\rho} \sum_{i=1}^n \mathcal{S} \left( \theta_i^{(m-1)} - \rho \cdot f^{(m)}(x_i), y_i \right)$ 
  for  $i \leftarrow 1, \dots, n$  do
     $\theta_i^{(m)} \leftarrow \theta_i^{(m-1)} - \eta \left( \rho^{(m)} \cdot f^{(m)}(x_i) \right)$ 
  end
end
    
```

A prediction $y|x$ on a new input x is made in the form of a conditional distribution P_θ , whose parameters θ are obtained by an additive combination of M base learner outputs (corresponding to the M gradient boosting stages) and an initial $\theta^{(0)}$. Note that θ can be a vector of parameters (not limited to be scalar valued), and they completely determine the probabilistic prediction $y|x$. For example, when using the Normal distribution, $\theta = (\mu, \log \sigma)$ in our experiments. To obtain the predicted parameter θ for some x , each of the base learners $f^{(m)}$ take x as their input. Here $f^{(m)}$ collectively refers to the set of base learners, one per parameter, of stage m . For example, for a Normal distribution with parameters μ and $\log \sigma$, there will be two base learners, $f_\mu^{(m)}$ and $f_{\log \sigma}^{(m)}$ per stage, collectively denoted as $f^{(m)} = (f_\mu^{(m)}, f_{\log \sigma}^{(m)})$. The predicted outputs are scaled with a stage-specific scaling factor $\rho^{(m)}$, and a common learning rate η :

$$y|x \sim P_\theta(x), \quad \theta = \theta^{(0)} - \eta \sum_{m=1}^M \rho^{(m)} \cdot f^{(m)}(x).$$

The scaling factor $\rho^{(m)}$ is a single scalar, even if the distribution has multiple parameters. The model is learnt sequentially, a set of base learners $f^{(m)}$ and a scaling factor

$\rho^{(m)}$ per stage. The learning algorithm starts by first estimating a common $\theta^{(0)}$ such that it minimizes the sum of the scoring rule \mathcal{S} over the response variables from all training examples, essentially fitting the marginal distribution of y . This becomes the initial predicted parameter $\theta^{(0)}$ for all examples.

In each iteration m , the algorithm calculates, for each example i , the natural gradients $g_i^{(m)}$ of the scoring rule \mathcal{S} with respect to the predicted parameters of that example up to that stage, $\theta_i^{(m-1)}$. Note that $g_i^{(m)}$ has the same dimension as θ . A set of base learners for that iteration $f^{(m)}$ are fit to predict the corresponding components of the natural gradients $g_i^{(m)}$ of each example x_i .

The output of the fitted base learner is the projection of the natural gradient on to the range of the base learner class. This projected gradient is then scaled by a scaling factor $\rho^{(m)}$ since local approximations might not hold true very far away from the current parameter position. The scaling factor is chosen to minimize the overall true scoring rule loss along the direction of the projected gradient in the form of a line search. In practice, we found that implementing this line search by successive halving of ρ (starting with $\rho = 1$) until the scaled gradient update results in a lower overall loss relative to the previous iteration works reasonably well and is easy to implement.

Once the scaling factor $\rho^{(m)}$ is determined, the predicted per-example parameters are updated to $\theta_i^{(m)}$ by adding to each $\theta_i^{(m-1)}$ the negative scaled projected gradient $\rho^{(m)} \cdot f^{(m)}(x_i)$ which is further scaled by a small learning rate η (typically 0.1 or 0.01).

The pseudo-code is presented in Algorithm 1. For very large datasets computational performance can be easily improved by simply randomly sub-sampling mini-batches within the fit() operation.

2.5. Qualitative Analysis and Discussion

Boosting for probabilistic prediction. Our boosting approach generalizes gradient boosting to predict conditional distributions. For instance, if the user specifies the conditional distribution to be a Normal distribution with a *fixed* variance and uses the log scoring rule, our approach recovers the standard boosting algorithm with MSE loss (modulo the per-leaf line search). The advantage of NGBoost is that users are also free to specify any other family of distributions identified by a set of real-valued parameters and allow all of those parameters to vary over the covariates, not just the mean. NGBoost thus trivially extends to a variety of use cases, such as negative binomial boosting (for counts), Gamma or Weibull boosting (for survival prediction, with or without right-censored data) etc.




Figure 4. Contrasting the learning dynamics between using the ordinary gradient (top row) vs. the natural gradient (bottom row) for the purpose of gradient boosting the parameters of a Normal distribution on a toy data set. With ordinary gradients, we observe that “lucky” examples that are accidentally close to the initial predicted mean dominate the learning. This is because, under the ordinary gradient, the variance of those examples that have the correct mean gets adjusted much more aggressively than the wrong means of the “unlucky” examples. This results in simultaneous overfitting of the “lucky” examples in the middle and underfitting of the “unlucky” examples at the ends. Under the natural gradient, all the updates are better balanced.

Multiparameter boosting. These wide-ranging extensions of gradient boosting are made possible by turning the distributional prediction problem into a problem of jointly estimating p functions of x , one per parameter, according to the scoring rule objective. In this setting, an overall line search for the stage multiplier (as opposed to per-leaf line search) is an inevitable consequence. However, our use of natural gradient makes this less of a problem as the gradients of all the examples come “optimally pre-scaled” (in both the relative magnitude between parameters, and across examples) due to the inverse Fisher Information factor. The use of ordinary gradients instead would be sub-optimal, as shown in Figure 4. With the natural gradient the parameters converge at approximately the same rate despite different conditional means, variances, and “distances” from the initial marginal distribution, even while being subjected to a common scaling factor $\rho^{(m)}$ in each iteration. We attribute this stability to the “optimal pre-scaling” property of the natural gradient.

Parametrization. When the probability distribution is in the exponential family and the parameters are the natural parameters of that family, then a Newton-Raphson step is equivalent to a natural gradient descent step. However, in other parametrizations and distributions, the equivalence need not hold. This is especially important in the boosting context because, depending on the inductive biases of the base learners, certain parametrization choices may result in more suitable model spaces than others. For example, one setting we are particularly interested in is the two-parameter Normal distribution. Though it is in the exponential family, we use a mean (μ) and log-scale ($\log \sigma$) parametrization for both ease of implementation and modeling convenience (to disentangle magnitude of predictions from uncertainty estimates). Since natural gradients are in-

variant to parametrization this does not pose a problem, whereas the Newton-Raphson method would fail as the problem is no longer convex in this parametrization.

3. Experiments

Our experiments use datasets from the UCI Machine Learning Repository, and follow the protocol first proposed in [Hernández-Lobato and Adams \(2015\)](#). For all datasets, we hold out a random 10% of the examples as a test set. From the other 90% we initially hold out 20% as a validation set to select M (the number of boosting stages) that gives the best log-likelihood, and then retrain on the entire 90% using the chosen M . The retrained model is then made to predict on the held-out 10% test set. This entire process is repeated 20 times for all datasets except Protein and Year MSD, for which it is repeated 5 times and 1 time respectively.

For all experiments, NGBoost was configured with the Normal distribution, decision tree base learner with a maximum depth of three levels, and log scoring rule. The Year MSD dataset, being extremely large relative to the rest, was fit using a learning rate η of 0.1 while the rest of the datasets were fit with a learning rate of 0.01. In general we recommend small learning rates, subject to computational feasibility. For the Year MSD dataset we use a mini-batch size of 10%, for all other datasets we use 100%.

Probabilistic regression. The quality of predictive uncertainty is captured in the average negative log-likelihood (NLL) (i.e. $\log \hat{P}_\theta(y|x)$) as measured on the test set.

Our comparison in this task is against other probabilistic prediction methods. We compare our results against MC dropout ([Gal and Ghahramani, 2016](#)), Deep Ensembles

Table 1. Comparison of probabilistic regression performance on regression benchmark UCI datasets. Results for MC dropout, Deep Ensembles, and Concrete Dropout are reported from Gal and Ghahramani (2016), Lakshminarayanan et al. (2017), and Gal et al. (2017a) respectively. NGBoost offers competitive performance in terms of NLL, especially on smaller datasets. The best method for each dataset is bolded, as are those with standard errors that overlap with the best method.

Dataset	N	NLL				
		NGBoost	MC dropout	Deep Ensembles	Concrete Dropout	Gaussian Process
Boston	506	2.43 \pm 0.15	2.46 \pm 0.25	2.41 \pm 0.25	2.72 \pm 0.01	2.37 \pm 0.24
Concrete	1030	3.04 \pm 0.17	3.04 \pm 0.09	3.06 \pm 0.18	3.51 \pm 0.00	3.03 \pm 0.11
Energy	768	0.60 \pm 0.45	1.99 \pm 0.09	1.38 \pm 0.22	2.30 \pm 0.00	0.66 \pm 0.17
Kin8nm	8192	-0.49 \pm 0.02	-0.95 \pm 0.03	-1.20 \pm 0.02	-0.65 \pm 0.00	-1.11 \pm 0.03
Naval	11934	-5.34 \pm 0.04	-3.80 \pm 0.05	-5.63 \pm 0.05	-5.87 \pm 0.05	-4.98 \pm 0.02
Power	9568	2.79 \pm 0.11	2.80 \pm 0.05	2.79 \pm 0.04	2.75 \pm 0.01	2.81 \pm 0.05
Protein	45730	2.81 \pm 0.03	2.89 \pm 0.01	2.83 \pm 0.02	2.81 \pm 0.00	2.89 \pm 0.02
Wine	1588	0.91 \pm 0.06	0.93 \pm 0.06	0.94 \pm 0.12	1.70 \pm 0.00	0.95 \pm 0.06
Yacht	308	0.20 \pm 0.26	1.55 \pm 0.12	1.18 \pm 0.21	1.75 \pm 0.00	0.10 \pm 0.26
Year MSD	515345	3.43 \pm NA	3.59 \pm NA	3.35 \pm NA	NA \pm NA	NA \pm NA

Table 2. Comparison of probabilistic regression performance on regression benchmark UCI datasets while ablating key components of NGBoost. Multiparameter boosting must be used in tandem with the natural gradient to increase performance. Bolding is as in Table 1.

Dataset	N	NLL		
		NGBoost	Multiparameter Boosting	Homoscedastic Boosting
Boston	506	2.43 \pm 0.15	3.17 \pm 0.13	2.79 \pm 0.42
Concrete	1030	3.04 \pm 0.17	3.94 \pm 0.09	3.22 \pm 0.29
Energy	768	0.60 \pm 0.45	3.24 \pm 0.09	0.68 \pm 0.25
Kin8nm	8192	-0.49 \pm 0.02	-0.52 \pm 0.03	-0.37 \pm 0.05
Naval	11934	-5.34 \pm 0.04	-3.46 \pm 0.00	-4.35 \pm 0.07
Power	9568	2.79 \pm 0.11	3.79 \pm 0.13	2.66 \pm 0.11
Protein	45730	2.81 \pm 0.03	3.04 \pm 0.02	2.86 \pm 0.01
Wine	1588	0.91 \pm 0.06	0.93 \pm 0.07	1.34 \pm 0.67
Yacht	308	0.20 \pm 0.26	3.29 \pm 0.20	2.02 \pm 0.21
Year MSD	515345	3.43 \pm NA	3.60 \pm NA	3.63 \pm NA

(Lakshminarayanan et al., 2017), Gaussian Process Regression (Rasmussen, 2003), and Concrete Dropout (Gal et al., 2017b). MC dropout fits a neural network to the dataset and interprets Bernoulli dropout as a variational approximation for Bayesian inference, obtaining predictive uncertainty by integrating over Monte Carlo samples. Deep Ensembles fit an ensemble of neural networks to the dataset and obtain predictive uncertainty by making an approximation to the Gaussian mixture arising out of the ensemble. Concrete Dropout improves upon MC dropout by employing a continuous relaxation of the Bernoulli distribution to automatically tune the dropout probability. Results for the Year MSD dataset are unavailable for Concrete Dropout because they were not reported by Gal et al. (2017a), and unavailable for Gaussian Processes because the necessary computations for gradient-based optimization of hyper-parameters did not fit in memory. Details of the parameter tuning for Gaussian Processes are in the appendix. Our probabilistic regression results are summarized in Table 1.

Ablation. We compare the NLL of our full NGBoost algorithm on these data vs. that of the following comparators: (a) NGBoost using the ordinary gradient (“Multiparameter Boosting”). To implement this, the gradient is not multiplied by the inverse of the Riemannian metric before fitting the weak learner. (b) NGBoost assuming a homoscedastic variance $\sigma^2(x) = \sigma^2$ (“Homoscedastic Boosting”). To do this we use NGBoost to estimate the expected value μ (with σ^2 fixed, equivalent to standard boosting), then estimate σ^2 as the empirical variance of the residuals on the training set. Note that the natural gradient plays no meaningful role when there is only a single parameter estimated with NGBoost. Our ablation results are summarized in Table 2.

Point estimation. Although NGBoost is not specifically designed for point estimation, it is easy to extract point estimates of expectations $\hat{\mathbb{E}}[y|x]$ from the estimated distributions $\hat{P}(y|x)$. We use this approach in a third evaluation to compare the same NGBoost models as above to the sklearn implementations of random forests, standard gradient boosting, and elastic net regression. Predictive per-

Table 3. Comparison of point-estimation performance on regression benchmark UCI datasets. Although not optimized for point estimation, NGBoost still offers competitive performance. Bolding is as in Table 1.

Dataset	N	RMSE			
		NGBoost	Elastic Net	Random Forest	Gradient Boosting
Boston	506	2.94 \pm 0.53	4.08 \pm 0.16	2.97 \pm 0.30	2.46 \pm 0.32
Concrete	1030	5.06 \pm 0.61	12.1 \pm 0.05	5.29 \pm 0.16	4.46 \pm 0.29
Energy	768	0.46 \pm 0.06	2.75 \pm 0.03	0.52 \pm 0.09	0.39 \pm 0.02
Kin8nm	8192	0.16 \pm 0.00	0.20 \pm 0.00	0.15 \pm 0.00	0.14 \pm 0.00
Naval	11934	0.00 \pm 0.00	0.00 \pm 0.00	0.00 \pm 0.00	0.00 \pm 0.00
Power	9568	3.79 \pm 0.18	4.42 \pm 0.00	3.26 \pm 0.03	3.01 \pm 0.10
Protein	45730	4.33 \pm 0.03	5.20 \pm 0.00	3.60 \pm 0.00	3.95 \pm 0.00
Wine	1588	0.63 \pm 0.04	0.58 \pm 0.00	0.50 \pm 0.01	0.53 \pm 0.02
Yacht	308	0.50 \pm 0.20	7.65 \pm 0.21	0.61 \pm 0.08	0.42 \pm 0.09
Year MSD	515345	8.94 \pm NA	9.49 \pm NA	9.05 \pm NA	8.73 \pm NA

formance in this evaluation is captured by the root mean squared-error (RMSE) of the predictions on the test set. We performed hyperparameter tuning for each of the comparator methods using the same validation procedure as described above, although optimizing for RMSE instead of NLL for all methods except NGBoost. The hyperparameter grids we used for tuning are available in the supplement. Our point prediction results are summarized in Table 3.

4. Conclusions

NGBoost is a method for probabilistic prediction with state-of-the-art performance on a variety of datasets. NGBoost combines a multiparameter boosting algorithm with the natural gradient to efficiently estimate how parameters of the presumed outcome distribution vary with the observed features.

NGBoost performs as well or better than existing methods for probabilistic regression, especially on smaller datasets, but also has other major advantages: NGBoost is flexible, scalable, and easy-to-use. We have not quantified these advantages in this paper (since they would be largely irrelevant without first establishing performance), but many of the benefits are self-evident. Unlike problem-specific approaches, NGBoost handles classification, regression, survival problems, etc. using the same software package and interface. NGBoost scales to large numbers of features or observations with the same favorable complexity of traditional boosting algorithms. And no expert knowledge of deep learning, Bayesian statistics, or Monte-Carlo methods is required to use NGBoost: it works out of the box.

Our ablation experiments demonstrate that multiparameter boosting and the natural gradient work together to improve performance. Assuming a uniform variance across all covariates works reasonably well for some datasets, as would be expected, but this is not always the case. However, using multiparameter boosting to relax the homoscedasticity

assumption most often results in even worse performance, likely due to poor training dynamics. NGBoost employs the natural gradient to correct the training dynamics of multiparameter boosting. The result is performance that is almost always better than assuming homoscedasticity, sometimes by a large margin.

Furthermore, the advantages of probabilistic regression come almost “for free”. On point estimation tasks NGBoost performs better than elastic net, about on par with random forests, and within striking distance of gradient boosting. This is despite the fact that the NGBoost models were a) optimized for NLL, not to minimize RMSE and b) less aggressively tuned. Thus, although point prediction will always be best with a dedicated model for that purpose, the loss in RMSE is not substantial if NGBoost is used in order to support probabilistic regression.

There are many avenues for future work. This paper is focused only on regression problems for clarity of exposition, but NGBoost is also applicable to classification (trivially) and to survival problems with right-censored data (using the censored likelihood as a scoring rule). NGBoost could also be used in *joint* prediction problems: by modeling two outcomes z and y with a jointly parameterized conditional distribution $P_{\theta}(z, y|x)$, a single NGBoost model could answer any question like “what is the probability that it rains more than 4 inches *and* the temperature is greater than 17C tomorrow?”.

Some further technical innovations are also worth exploring. The natural gradient loses its invariance property with finite step sizes, which we can address with differential equation solvers for higher-order invariance (Song et al., 2018). Better tree-based base learners and regularization (e.g. Chen and Guestrin (2016); Ke et al. (2017)) are also likely to improve performance.

ACKNOWLEDGEMENTS

This work was funded in part by the National Institutes of Health.

References

- Amari, S.-i. (1998). Natural Gradient Works Efficiently in Learning. *Neural Computation*, page 29.
- Avati, A., Duan, T., Jung, K., Shah, N. H., and Ng, A. (2019). Countdown Regression: Sharp and Calibrated Survival Predictions. In *Uncertainty in Artificial Intelligence*. arXiv: 1806.08324.
- Avati, A., Jung, K., Harman, S., Downing, L., Ng, A., and Shah, N. H. (2018). Improving palliative care with deep learning. *BMC Medical Informatics and Decision Making*, 18(4):122.
- Chen, T. and Guestrin, C. (2016). XGBoost: A Scalable Tree Boosting System. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '16, pages 785–794, New York, NY, USA. ACM. event-place: San Francisco, California, USA.
- Chipman, H. A., George, E. I., and McCulloch, R. E. (2010). BART: Bayesian additive regression trees. *The Annals of Applied Statistics*, 4(1):266–298.
- Dawid, A. P. (2007). The geometry of proper scoring rules. *Annals of the Institute of Statistical Mathematics*, 59(1):77–93.
- Dawid, A. P. and Musio, M. (2014). Theory and Applications of Proper Scoring Rules. *METRON*, 72(2):169–183. arXiv: 1401.0398.
- Friedman, J. H. (2001). Greedy Function Approximation: A Gradient Boosting Machine. *The Annals of Statistics*, 29(5):1189–1232.
- Gal, Y. and Ghahramani, Z. (2016). Dropout As a Bayesian Approximation: Representing Model Uncertainty in Deep Learning. In *International Conference on Machine Learning*, ICML'16, pages 1050–1059. JMLR.org. event-place: New York, NY, USA.
- Gal, Y., Hron, J., and Kendall, A. (2017a). Concrete Dropout. In Guyon, I., Luxburg, U. V., Bengio, S., Wallach, H., Fergus, R., Vishwanathan, S., and Garnett, R., editors, *Advances in Neural Information Processing Systems 30*, pages 3581–3590. Curran Associates, Inc.
- Gal, Y., Hron, J., and Kendall, A. (2017b). Concrete dropout. In *Advances in Neural Information Processing Systems*, pages 3581–3590.
- Gebetsberger, M., Messner, J. W., Mayr, G. J., and Zeileis, A. (2018). Estimation Methods for Nonhomogeneous Regression Models: Minimum Continuous Ranked Probability Score versus Maximum Likelihood. *Monthly Weather Review*, 146(12):4323–4338.
- Gneiting, T. and Katzfuss, M. (2014). Probabilistic Forecasting. *Annual Review of Statistics and Its Application*, 1(1):125–151.
- Gneiting, T. and Raftery, A. E. (2005). Weather Forecasting with Ensemble Methods. *Science*, 310(5746):248–249.
- Gneiting, T. and Raftery, A. E. (2007). Strictly Proper Scoring Rules, Prediction, and Estimation. *Journal of the American Statistical Association*, 102(477):359–378.
- Hernández-Lobato, J. M. and Adams, R. P. (2015). Probabilistic Backpropagation for Scalable Learning of Bayesian Neural Networks. In *International Conference on Machine Learning*, ICML'15, pages 1861–1869. JMLR.org. event-place: Lille, France.
- Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., and Liu, T.-Y. (2017). LightGBM: A Highly Efficient Gradient Boosting Decision Tree. In Guyon, I., Luxburg, U. V., Bengio, S., Wallach, H., Fergus, R., Vishwanathan, S., and Garnett, R., editors, *Advances in Neural Information Processing Systems 30*, pages 3146–3154. Curran Associates, Inc.
- Kruchten, N. (2016). Machine learning meets economics.
- Lakshminarayanan, B., Pritzel, A., and Blundell, C. (2017). Simple and Scalable Predictive Uncertainty Estimation using Deep Ensembles. In Guyon, I., Luxburg, U. V., Bengio, S., Wallach, H., Fergus, R., Vishwanathan, S., and Garnett, R., editors, *Advances in Neural Information Processing Systems 30*, pages 6402–6413. Curran Associates, Inc.
- Machete, R. L. (2013). Contrasting probabilistic scoring rules. *Journal of Statistical Planning and Inference*, 143(10):1781–1790.
- Martens, J. (2014). New insights and perspectives on the natural gradient method. Technical report. arXiv: 1412.1193.
- Neal, R. M. (1996). *Bayesian Learning for Neural Networks*. Springer-Verlag, Berlin, Heidelberg.
- Rasmussen, C. E. (2003). Gaussian processes in machine learning. In *Summer School on Machine Learning*, pages 63–71. Springer.
- Rasmussen, C. E. and Williams, C. K. I. (2005). *Gaussian Processes for Machine Learning (Adaptive Computation and Machine Learning)*. The MIT Press.

- Song, Y., Song, J., and Ermon, S. (2018). Accelerating Natural Gradient with Higher-Order Invariance. In *International Conference on Machine Learning*, pages 4713–4722.
- Stasinopoulos, D. M., Rigby, R. A., et al. (2007). Generalized additive models for location scale and shape (gamlss) in r. *Journal of Statistical Software*, 23(7):1–46.